

Lecture 8: Reinforcement Learning and Generative Models

Isaiah Hull^{1,2}

¹BI Norwegian Business School

²CogniFrame

December 5, 2023



Unsupervised Learning

Lecture 8: Overview

1. Generative Models.
2. Reinforcement Learning.
3. Applications.

1. Generative Models

Generative Models

Types of ML Models

- ▶ Discriminative models:
 - ▶ Applied to classification or regression.
 - ▶ Input: set of features.
 - ▶ Output: probabilities or predicted values.

Generative Models

Types of ML Models

- ▶ Generative models:
 - ▶ Learn the data distribution.
 - ▶ Produce new class examples.

Generative Models

Advancements in Generative ML

- ▶ Significant progress in generative ML.
- ▶ Notably:
 - ▶ Variational autoencoders (VAEs).
 - ▶ Generative adversarial networks (GANs).
- ▶ Breakthroughs in image, text, and music generation.

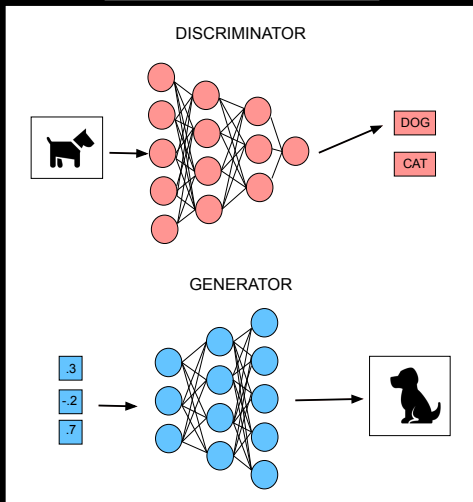
Generative Models

Generative ML in Economics

- ▶ Limited use of generative models in economics and finance.
- ▶ Growing interest in GANs.
- ▶ Example applications:
 - ▶ Athey et al. (2019).
 - ▶ Kaji et al. (2018).
- ▶ Potential future uses in the field.

Generative Models

Types of ML Models



Source: Hull (2021).

Generative Models

Variational Autoencoders

- ▶ Autoencoders:
 - ▶ Two networks: encoder and decoder with shared weights.
 - ▶ Encoder: Transforms input to latent state.
 - ▶ Decoder: Reconstructs features from latent state.
 - ▶ Trained via reconstruction loss.
- ▶ Uses: dimensionality reduction, generative tasks (images, music, texts).

Generative Models

Issues with Autoencoders

- ▶ Location and distribution of latent states:
 - ▶ Latent states are in R^N .
 - ▶ Clustering is not explicitly defined.
 - ▶ Affects determination of valid latent states for generation.
- ▶ Performance for unobserved latent states:
 - ▶ Only trained for specific examples.
 - ▶ Perturbations might lead to unconvincing outputs.

Generative Models

Variational Autoencoders (VAEs)

- ▶ Developed to address autoencoder limitations.
- ▶ Components:
 - ▶ Mean layer.
 - ▶ Log variance layer.
 - ▶ Sampling layer (from normal distribution).
- ▶ Output of sampling is the latent state.
- ▶ Encoder yields different latent states for same inputs.

Generative Models

VAE Architecture and Loss

- ▶ VAE modifies loss to include Kullback-Leibler (KL) divergence.
- ▶ KL divergence:
 - ▶ Penalizes distance between normal distributions.
 - ▶ Targets a mean and log variance of zero.

Generative Models

VAE Architecture and Loss

- ▶ Eliminates determinism of latent states.
- ▶ Improves generative performance.
- ▶ Solves sampling problem with sampling layer.
- ▶ Adjusts latent distribution with KL divergence.

Generative Models

VAEs in TensorFlow

- ▶ Detailed exploration: See Kingma and Welling (2019).
- ▶ Data used: GDP growth from 1961:Q2 to 2020:Q1 for 25 OECD countries.
- ▶ Previous usage: Dimensionality reduction.
- ▶ Objective: Train VAE to generate similar series.

Generative Models

Prepare GDP Growth Data for Use in a VAE

```
import tensorflow as tf
import pandas as pd
import numpy as np
```

```
# Load and transpose data.
```

```
GDP = pd.read_csv(data_path+'gdp_growth.csv',
                  index_col = 'Date').T
```

```
# Print data preview.
```

```
print(GDP.head())
```

Generative Models

Model Architecture Summary

Time	4/1/61	7/1/61	10/1/61	1/1/62
AUS	-1.097616	-0.715607	1.139175	2.806800
AUT	-0.349959	1.256452	0.227988	1.463310
BEL	1.167163	1.275744	1.381074	1.346942
CAN	2.529317	2.409293	1.396820	2.650176
CHE	1.355571	1.242126	1.958044	0.575396

Generative Models

Model Architecture Summary

Convert data to numpy array.

GDP = np.array(GDP)

Set number of countries and quarters.

nCountries, nQuarters = GDP.shape

Set number of latent nodes and batch size.

latentNodes = 2

batchSize = 1

Generative Models

VAE Model Architecture

- ▶ Contains encoder and decoder.
- ▶ Differs from autoencoder.
- ▶ Latent states sampled from independent normal distributions.

Generative Models

Sampling Layer in VAE

- ▶ Takes two parameters as inputs.
- ▶ Draws epsilon from standard normal distribution for each output node.
- ▶ Transforms each draw using mean and lvar parameters.

Generative Models

Model Architecture Summary

Define function for sampling layer.

```
def sampling(params, batchSize = batchSize, latentNodes = latentNodes):  
    mean, lvar = params  
    epsilon = tf.random.normal(shape=(batchSize, latentNodes))  
    return mean + tf.exp(lvar / 2.0) * epsilon
```

Generative Models

Encoder Model in VAE

- ▶ Takes full time series for a country as input.
- ▶ Mean and log variance layers differ from autoencoder.
- ▶ Mean and log variances layers have nodes for latent state.

Generative Models

Define Encoder Model for VAE

Define input layer for encoder.

```
encoderInput = tf.keras.layers.Input(shape = (nQuarters))
```

Define latent state.

```
latent = tf.keras.layers.Input(shape = (latentNodes))
```

Define mean layer.

```
mean = tf.keras.layers.Dense(latentNodes)(encoderInput)
```

Generative Models

Define Encoder Model for VAE

Define log variance layer.

```
lvar = tf.keras.layers.Dense(latentNodes)(encoderInput)
```

Define sampling layer.

```
encoded = tf.keras.layers.Lambda(sampling,  
                                output_shape=(latentNodes,))([mean, lvar])
```

Define model for encoder.

```
encoder = tf.keras.Model(encoderInput, [mean, lvar, encoded])
```

Generative Models

Further Details on Encoder

- ▶ Mean and lvar layers parameterize normal distributions.
- ▶ Lambda layer defined with the sampling function.
- ▶ Lambda layer takes mean and lvar parameters.

Generative Models

Further Details on Encoder

- ▶ Encoder model defined functionally.
- ▶ Inputs: quarterly GDP growth observations.
- ▶ Outputs: mean layer, log variance layer, and sampled outputs.

Generative Models

Functional Models for Decoder and VAE

- ▶ Decoder: accepts latent state and produces a reconstruction.
- ▶ VAE: takes a time series and transforms it into its reconstruction.

Generative Models

Defining the Loss Function

- ▶ Two-part loss function:
 - ▶ Reconstruction loss: same as autoencoder.
 - ▶ KL divergence: measures distance of sampling layer distributions from a standard normal.
- ▶ Further from standard normal distribution implies higher penalty.

Generative Models

Define Decoder Model for VAE

Define output for decoder.

```
decoded = tf.keras.layers.Dense(nQuarters, activation = 'linear')(latent)
```

Define the decoder model.

```
decoder = tf.keras.Model(latent, decoded)
```

Define functional model for autoencoder.

```
vae = tf.keras.Model(encoderInput, decoder(encoded))
```

Generative Models

Define VAE Loss

Compute the reconstruction component of the loss.

```
reconstruction = tf.keras.losses.binary_crossentropy(vae.inputs[0], vae.outputs[0])
```

Compute the KL loss component.

```
kl = -0.5 * tf.reduce_mean(1 + lvar - tf.square(mean) - tf.exp(lvar), axis = -1)
```

Combine the losses and add them to the model.

```
combinedLoss = reconstruction + kl  
vae.add_loss(combinedLoss)
```

Generative Models

Trained Variational Autoencoder

- ▶ Capabilities of the trained VAE:
 - ▶ Use `predict()` method of `vae` to reconstruct given time series.
 - ▶ Generate realization of latent state for specific inputs, e.g., GDP growth for the US.
 - ▶ Perturb latent states by adding noise and generate new time series using `predict()` of decoder.

Generative Models

Compile and Fit VAE

Compile the model.

```
vae.compile(optimizer='adam')
```

Fit model.

```
vae.fit(GDP, batch_size = batchSize, epochs = 100)
```

Generative Models

Generate Latent States and Time Series with Trained VAE

Generate series reconstruction.

```
prediction = vae.predict(GDP[0,:].reshape(1,236))
```

Generate (random) latent state from inputs.

```
latentState = encoder.predict(GDP[0,:].reshape(1,236))
```

Perturb latent state.

```
latentState[0] = latentState[0] + np.random.normal(1)
```

Pass perturbed latent state to decoder.

```
decoder.predict(latentState)
```

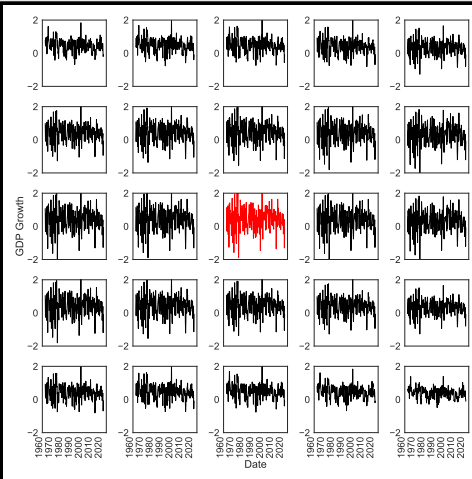

Generative Models

Generated Time Series for U.S. GDP Growth

- ▶ Based on a latent state realization for the U.S. GDP growth series.
- ▶ Perturbations over a 5x5 grid:
 - ▶ Rows: Add evenly-spaced values over the $[-1,1]$ interval to the first latent state.
 - ▶ Columns: Add evenly-spaced values over the $[-1,1]$ interval to the second latent state.
- ▶ Center series (in red): Adds $[0,0]$ - Represents the original latent state.

Generative Models

VAE-Generated Time Series for U.S. GDP Growth



Source: Hull (2021).

Generative Models

VAE Application and Flexibility

- ▶ VAE architecture's broad applicability:
 - ▶ Add convolutional layers to encoder/decoder \Rightarrow Generate images.
 - ▶ Incorporate LSTM cells \Rightarrow generate text or music.
- ▶ LSTM-based architecture potential:
 - ▶ Possible improvements in time series generation.
 - ▶ A shift from the dense network approach used in the given example.

Generative Models

Generative Adversarial Networks

- ▶ Dominant generative ML models:
 - ▶ Variational Autoencoders (VAEs).
 - ▶ Generative Adversarial Networks (GANs).
- ▶ VAEs: Manipulate latent states for granular control.
- ▶ GANs: Convincing generation of class examples.
- ▶ Notable GAN results: Highly convincing generated images.

Generative Models

GANs vs VAEs

- ▶ VAEs:
 - ▶ Encoder + Decoder, joined by sampling layer.
- ▶ GANs:
 - ▶ Generator + Discriminator.
 - ▶ Generator: Latent state input \Rightarrow class example (e.g. GDP growth time series).
 - ▶ Discriminator: Differentiate between real and fake examples.

Generative Models

GANs vs VAEs

- ▶ Adversarial Network:
 - ▶ Combines Generator + Discriminator.
 - ▶ Shares weights with both networks.
 - ▶ Trains the generator to maximize discriminator's loss.

Generative Models

GAN: Zero Sum Game

- ▶ Goodfellow et al. (2017):
 - ▶ Two networks in a zero-sum game.
 - ▶ Discriminator: $v(g,d)$.
 - ▶ Generator: $-v(g,d)$.
 - ▶ Evolutionary equilibrium.

Generative Models

GAN: Zero Sum Game

- ▶ Equilibrium Equation: $g^* = \arg \min_g \max_d v(g, d)$
- ▶ Freeze discriminator weights during adversarial training.
- ▶ Improve generation, rather than weakening the discriminator.

Generative Models

GAN Architecture Illustrated

- ▶ GAN's generator and discriminator networks.
- ▶ Generator produces novel examples.
- ▶ Discriminator combines novel and true examples for classification.
- ▶ Adversarial Network:
 - ▶ Trains generator.
 - ▶ Attached to discriminator with frozen weights.
 - ▶ Iterative training for equilibrium.


Generative Models

Practical Application: GDP Growth Data

- ▶ GDP growth data (similar to VAE section).
- ▶ Train GAN for credible GDP growth time series generation from random vector input.
- ▶ Approach to GAN construction from Krohn et al. (2020).

Generative Models

Depiction of the Generator and Discriminator from a GAN.



Source: Hull (2021).

Generative Models

Prepare GDP Growth Data for Use in a GAN.

```
import tensorflow as tf
import pandas as pd
import numpy as np
```

```
# Load and transpose data.
```

```
GDP = pd.read_csv(data_path+'gdp_growth.csv',
                  index_col = 'Date').T
```

```
# Convert pandas DataFrame to numpy array.
```

```
GDP = np.array(GDP)
```

Generative Models

Generative Model Definition

- ▶ Simple VAE model.
- ▶ Generator's input:
 - ▶ Vector with two elements.
 - ▶ Analogous to latent vector in VAE.

Generative Models

Generative Model Definition

- ▶ Generator viewed as a decoder.
- ▶ Architecture:
 - ▶ Start with bottleneck-type layer.
 - ▶ Upsample to produce GDP growth time series.

Generative Models

Generator's Architecture

- ▶ Simplest version:
 - ▶ Input layer: Accepts latent vector.
 - ▶ Output layer: Upsamples input layer.
- ▶ Output layer:
 - ▶ Represents GDP growth values.
 - ▶ Activation function: Linear.

Generative Models

Define the Generative Model of a GAN

Set dimension of latent state vector.

nLatent = 2

Set number of countries and quarters.

nCountries, nQuarters = GDP.shape

Define input layer.

generatorInput = tf.keras.layers.Input(shape = (nLatent,))

Generative Models

Define the Generative Model of a GAN

Define hidden layer.

```
generatorHidden = tf.keras.layers.Dense(16, activation='relu')(generatorInput)
```

Define generator output layer.

```
generatorOutput = tf.keras.layers.Dense(236, activation='linear')(generatorHidden)
```

Define generator model.

```
generator = tf.keras.Model(inputs = generatorInput, outputs = generatorOutput)
```

Generative Models

Discriminator Definition

- ▶ Inputs:
 - ▶ Real GDP growth series.
 - ▶ Generated GDP growth series.
- ▶ Each series length: n_{Quarters} .
- ▶ Output: Probability of being a real GDP growth series for each input series.

Generative Models

Compilation Details

- ▶ Generator: Not compiled.
- ▶ Discriminator: Compiled.
- ▶ Use of an adversarial network to train the generator.

Generative Models

Define and Compile the Discriminator Model of a GAN

Define input layer.

```
discriminatorInput = tf.keras.layers.Input(shape = (nQuarters,))
```

Define hidden layer.

```
discriminatorHidden = tf.keras.layers.Dense(16,  
      activation='relu')(discriminatorInput)
```

Generative Models

Define and Compile the Discriminator Model of a GAN

Define discriminator output layer.

```
discriminatorOutput = tf.keras.layers.Dense(1,  
      activation='sigmoid')(discriminatorHidden)
```

Define discriminator model.

```
discriminator = tf.keras.Model(inputs = discriminatorInput,  
      outputs = discriminatorOutput)
```

Compile discriminator.

```
discriminator.compile(loss='binary_crossentropy',  
      optimizer=tf.optimizers.Adam(0.0001))
```

Generative Models

Adversarial Model

- ▶ Shares weights with the generator.
- ▶ Uses a frozen version of discriminator's weights.
- ▶ Discriminator weights:
 - ▶ Do not update during adversarial training.
 - ▶ Update during discriminator training.

Generative Models

Adversarial Network

- ▶ Input: Latent vector (same size as generator input).
- ▶ Output of generator: timeSeries (fake GDP growth time series).
- ▶ Set discriminator's trainability to 'False' during adversarial training.
- ▶ Network's output: Discriminator's output.
- ▶ Defined and compiled functional model: adversarial.

Generative Models

Training Process

- ▶ Train the discriminator.
- ▶ Train the adversarial network.

Generative Models

Define and Compile the Adversarial Model of a GAN

Define input layer for adversarial network.

```
adversarialInput = tf.keras.layers.Input(shape=(nLatent))
```

Define generator output as generated time series.

```
timeSeries = generator(adversarialInput)
```

Set discriminator to be untrainable.

```
discriminator.trainable = False
```

Generative Models

Define and Compile the Adversarial Model of a GAN

Set discriminator to be untrainable.

```
discriminator.trainable = False
```

Compute predictions from discriminator.

```
adversarialOutput = discriminator(timeSeries)
```

Define adversarial model.

```
adversarial = tf.keras.Model(adversarialInput, adversarialOutput)
```

Compile adversarial network.

```
adversarial.compile(loss='binary_crossentropy',  
                    optimizer=tf.optimizers.Adam(0.0001))
```

Generative Models

Train the Discriminator and the Adversarial Network

Set batch size.

batch, halfBatch = 12, 6

for j **in** range(1000):

Draw real training data.

idx = np.random.randint(nCountries,
size = halfBatch)

real_gdp_series = GDP[idx, :]

Generate fake training data.

latentState = np.random.normal(size=[halfBatch, nLatent])

fake_gdp_series = generator.predict(latentState)

Generative Models

Train the Discriminator and the Adversarial Network

Combine input data.

```
features = np.concatenate((real_gdp_series,  
                           fake_gdp_series))
```

Create labels.

```
labels = np.ones([batch,1])  
labels[halfBatch:, :] = 0
```

Generative Models

Train the Discriminator and the Adversarial Network

Train discriminator.

```
discriminator.train_on_batch(features, labels)
```

Generate latent state for adversarial net.

```
latentState = np.random.normal(size=[batch, nLatent])
```

Generate labels for adversarial network.

```
labels = np.ones([batch, 1])
```

Train adversarial network.

```
adversarial.train_on_batch(latentState, labels)
```

Generative Models

Initialization

- ▶ Define the batch size.
- ▶ Begin training loop consisting of multiple steps.

Generative Models

Training the Discriminator

- ▶ Draw random integers for row selection in GDP matrix.
- ▶ Each row contains GDP growth time series (real samples for discriminator).
- ▶ Generate fake data:
 - ▶ Draw latent vectors.
 - ▶ Pass through generator.
- ▶ Combine real and fake series with labels (1 = real, 0 = fake).
- ▶ Train discriminator with combined data.

Generative Models

Training the Adversarial Network

- ▶ Generate a batch of latent states.
- ▶ Input latent states into generator.
- ▶ Train with intent to trick discriminator into classifying as real.
- ▶ Iterative training over two models.
- ▶ Stopping criteria: stable evolutionary equilibrium.


Generative Models

Model Evaluation

- ▶ Observe model losses over time.
- ▶ Approx. 500 iterations: neither model shows significant improvement.
- ▶ Indicates a stable evolutionary equilibrium reached.

Generative Models

Discriminator and Adversarial Model Losses by Training Iteration



Source: Hull (2021).

Generative Models

Adversarial Network Performance

- ▶ White noise vector inputs used.
- ▶ Adversarial network informed by discriminator's performance.
- ▶ Achieved fairly credible fake GDP growth series after 1000 training iterations.


Generative Models

Performance Enhancement

- ▶ More latent features could improve performance.
- ▶ Can use a more advanced model architecture, such as LSTM.

Generative Models

Example Fake GDP Growth Series.



Source: Hull (2021).

Generative Models

Applications in Economics and Finance

- ▶ Concentrated on generating simulated series.
- ▶ Alternative to more familiar methods in Monte Carlo simulation.
- ▶ Must generate realistic series and capture interdependencies.

Generative Models

Early Applications of GANs in Economics

- ▶ Athey et al. (2019) used Wasserstein GANs.
- ▶ Simulate data appearing similar to small existing datasets.
- ▶ Bypasses drawing randomly or generating simulated series with deficiencies.

Generative Models

Value of GANs in Economics

- ▶ Athey et al. evaluated estimators using WGAN generated data.
- ▶ Kaji et al. (2018) used WGANs for indirect inference.
- ▶ Indirect inference: estimating structural models in economics and finance.

Generative Models

Kaji et al.'s Approach

- ▶ Estimate a model for workers choosing wage and location.
- ▶ Parameters are structural.
- ▶ Coupled model simulation with a discriminator.

Generative Models

Further Applications

- ▶ GANs and VAEs in image and text generation.
- ▶ Visual counterfactual simulations with economic data.
- ▶ NLP in economics: text generation to study company press releases.

2. Deep Reinforcement Learning

Reinforcement Learning

Deep Reinforcement Learning

- ▶ In standard models, agents are rational optimizers.
- ▶ Agents form unbiased expectations about the future.
- ▶ Optimizers choose the exact optimum.
- ▶ No heuristics or rule-of-thumb used.

Reinforcement Learning

Deviation from Rational Optimizer Framework

- ▶ Palmer (2015).
- ▶ Policy rule formation, rather than rationality.
- ▶ Improved computational tractability by breaking requirements.

Reinforcement Learning

Reinforcement Learning in Economics

- ▶ Alternative to standard model.
- ▶ Described in Sutton and Barto (1998).
- ▶ Value discussed in Athey and Imbens (2019), Palmer (2015).
- ▶ Application: Hull (2015) for dynamic programming problems.

Reinforcement Learning

Agents in Reinforcement Learning

- ▶ Agents perform optimization.
- ▶ Limited information about the system state.
- ▶ Tradeoff: "exploration" vs "exploitation."
- ▶ Learning about the system vs optimizing known system parts.

Reinforcement Learning

Deep Q-learning Introduction

- ▶ Variant of reinforcement learning.
- ▶ Combines deep learning and reinforcement learning.
- ▶ Can solve high-dimensional state space problems.
- ▶ Solves rational optimizer's problem via deep Q-learning.

Reinforcement Learning

Dynamic Programming vs Q-learning

- ▶ Dynamic programming uses "look-up table" for value states.
- ▶ Iterative table updates until convergence.
- ▶ Solution found in the value function table.
- ▶ Q-learning constructs a state-action table.
- ▶ Example: State - capital stock; Action - level of consumption.

Reinforcement Learning

Q-table Update Mechanism

- ▶ Updated using temporal difference learning.
- ▶ Updates state-action pair (s_t, a_t) in iteration $i + 1$.
- ▶ Uses value in iteration i and adds learning rate.
- ▶ Multiplied by expected change in value from optimal action.

Reinforcement Learning

Updating the Q Table

$$Q_{(i+1)}(s_t, a_t) \leftarrow Q_i(s_t, a_t) + \lambda \left[r_t + \beta \max_a Q(k_{(t+1)}, a) - Q_i(s_t, a_t) \right] \quad (1)$$

Reinforcement Learning

Deep Q-learning

- ▶ Replaces look-up table with a deep neural network: “deep Q-network.”
- ▶ Approach introduced in Mnih et al. (2015).
- ▶ Originally used to train Q-networks for superhuman video game performance.

Reinforcement Learning

Applying Deep Q-learning to Economic Models

- ▶ Neoclassical business cycle model.
- ▶ Solutions implemented in TensorFlow.
- ▶ Two common TensorFlow options:
 - ▶ tf-agents: Native TensorFlow implementation.
 - ▶ keras-rl2: Uses high-level Keras API.
- ▶ We use keras-rl2 for simpler, familiar syntax.

Reinforcement Learning

Implementation Details

- ▶ Install keras-rl2 module.
- ▶ Submodules from rl module:
 - ▶ DQNAgent: Define a deep Q-learning agent.
 - ▶ EpsGreedyQPolicy: Set policy decisions process on training path.
 - ▶ SequentialMemory: Retain decision paths/outcomes for training.
- ▶ Import gym to define the model environment.

Reinforcement Learning

Install and Import Modules to Perform Deep Q-learning

Install keras-rl2.

```
pip install keras-rl2
```

Import numpy and tensorflow.

```
import numpy as np
```

```
import tensorflow as tf
```

Reinforcement Learning

Install and Import Modules to Perform Deep Q-learning

Import reinforcement learning modules from keras.

from rl.agents.dqn **import** DQNAgent

from rl.policy **import** EpsGreedyQPolicy

from rl.memory **import** SequentialMemory

Import module for comparing RL algorithms.

import gym

Reinforcement Learning

Setting up the Environment

- ▶ Set the number of capital nodes.
- ▶ Define environment: `planner`, a subclass of `gym.Env`.
- ▶ Specifies details of the social planner's reinforcement learning problem.

Reinforcement Learning

The planner Class: Initialization

- ▶ Define a discrete capital grid.
- ▶ Define action and observation spaces.
- ▶ Initialize the number of decisions to zero.
- ▶ Set the max number of decisions.
- ▶ Set node index of initial value of capital: 500 out of 1000.
- ▶ Set the production function parameter (α).

Reinforcement Learning

Action and Observation Spaces

- ▶ Both are discrete objects with 1000 nodes (defined using gym.spaces).
- ▶ Observation space: Entire state space (all capital nodes).
- ▶ Action space: Identical to the observation space.

Reinforcement Learning

Define Custom Reinforcement Learning Environment

Define number of capital nodes.

```
n_capital = 1000
```

Define environment.

```
class planner(gym.Env):  
    def __init__(self):  
        self.k = np.linspace(0.01, 1.0, n_capital)  
        self.action_space = \  
            gym.spaces.Discrete(n_capital)  
        self.observation_space = \  
            gym.spaces.Discrete(n_capital)  
        self.decision_count = 0
```

Reinforcement Learning

Define Custom Reinforcement Learning Environment

```
self.decision_max = 100
self.observation = 500
self.alpha = 0.33
def step(self, action):
    assert self.action_space.contains(action)
    self.decision_count += 1
    done = False
```

Reinforcement Learning

Define Custom Reinforcement Learning Environment

```
if(self.observation**self.alpha - action) > 0:
    reward = \
np.log(self.k[self.observation]**self.alpha -
self.k[action])
else:
    reward = -1000
self.observation = action
```

Reinforcement Learning

Define Custom Reinforcement Learning Environment

```
if (self.decision_count >= self.decision_max)\  
or reward == -1000:  
    done = True  
return self.observation, reward, done,\  
    {"decisions": self.decision_count}  
def reset(self):  
    self.decision_count = 0  
    self.observation = 500  
    return self.observation
```

Reinforcement Learning

Defining the step Method

- ▶ Method returns observation (state), reward (utility), reset indicator (done), and debugging info.
- ▶ Increment decision_count attribute.
- ▶ Initially sets done to False.
- ▶ Evaluates validity of agent's decision: selects positive consumption value.

Reinforcement Learning

Conditions for reset() Method

- ▶ Agent makes more than `decision_max` decisions.
- ▶ Agent chooses a non-positive consumption value.
- ▶ The `reset()` method reinitializes state and decision count.

Reinforcement Learning

Instantiating planner and Neural Network

- ▶ Instantiate a planner environment.
- ▶ Define a neural network using TensorFlow.
- ▶ Use the Sequential model.
- ▶ One dense layer with relu activation.
- ▶ Output layer should have `n_capital` nodes.

Reinforcement Learning

Define Custom Reinforcement Learning Environment

Instantiate planner environment.

```
env = planner()
```

Define model in TensorFlow.

```
model = tf.keras.models.Sequential()
```

```
model.add(tf.keras.layers.Flatten(input_shape=(1,) + env.observation_space.shape))
```

```
model.add(tf.keras.layers.Dense(32, activation='relu'))
```

```
model.add(tf.keras.layers.Dense(n_capital, activation='linear'))
```

Reinforcement Learning

Hyperparameters and Training

- ▶ Define environment and network.
- ▶ Use SequentialMemory for a "replay buffer" of 50,000 decision paths.
- ▶ Model uses epsilon-greedy policy:
 - ▶ $\epsilon = 0.30$.
 - ▶ Maximize utility 70% of the time.
 - ▶ Random decision for exploration 30% of the time.

Reinforcement Learning

DQNAgent Model Training

- ▶ Set hyperparameters of the DQNAgent model.
- ▶ Compile the model.
- ▶ Perform training.

Reinforcement Learning

Set Model Hyperparameters and Train.

Specify replay buffer.

```
memory = SequentialMemory(limit=10000, window_length=1)
```

Define policy used to make training-time decisions.

```
policy = EpsGreedyQPolicy(0.30)
```

Reinforcement Learning

Set Model Hyperparameters and Train.

Define deep Q-learning network (DQN).

```
dqn = DQNAgent(model=model, nb_actions=n_capital,  
               memory=memory, nb_steps_warmup=100,  
               gamma=0.95, target_model_update=1e-2,  
               policy=policy)
```

Compile and train model.

```
dqn.compile(tf.keras.optimizers.Adam(0.005), metrics=['mse'])  
dqn.fit(env, nb_steps=10000)
```

Reinforcement Learning

Monitoring the Training Process

- ▶ Two main observations from the training process:
 - ▶ Number of decisions per session increases across iterations.
 - ▶ Implication: Agent learns not to sharply draw down capital.
- ▶ Loss declines while the average reward starts to increase.
- ▶ Agent approaches optimality.

Reinforcement Learning

Summary

- ▶ Presents an alternative to standard computational economic methods.
- ▶ Deep Q-learning networks (DQN) in TensorFlow:
 - ▶ Can solve higher dimensional models.
 - ▶ Works in non-linear settings.
- ▶ Benefits:
 - ▶ No need to change model assumptions.
 - ▶ Minimizes numerical error introduction.

3. Applications

Applications

Colab Tutorial

- ▶ Generative Models
- ▶ Theoretical Models

References I

Athey, S. and G.W. Imbens (2019) “Machine Learning Methods Economist Should Know About,” *Annual Review of Economics*, 11, 685–725.

Athey, S., G.W. Imbens, J. Metzger, and E. Munro (2019) “Using Wasserstein Generative Adversarial Networks for the Design of Monte Carlo Simulations,” Working Paper 3824.

Goodfellow, I., Y. Bengio, and A. Courville (2017) *Deep Learning*, Cambridge, MA: MIT Press.

Hull, Isaiah (2021) *Machine Learning for Economics and Finance in TensorFlow 2*: Apress, 10.1007/978-1-4842-6373-0.

Kaji, T., E. Manresa, and G. Pouliot (2018) “Deep Inference: Artificial Intelligence for Structural Estimation,” working paper.

References II

Kingma, D.P. and M. Welling (2019) “An Introduction to Variational Autoencoders,” *Foundations and Trends in Machine Learning*, 12 (4), 307–392.

Krohn, J., G. Beyleveld, and A. Bassens (2020) *Deep Learning Illustrated: A Visual, Interactive Guide to Artificial Intelligence*: Addison-Wesley.

Palmer, N.M. (2015) *Individual and Social Learning: An Implementation of Bounded Rationality from First Principles* Ph.D. dissertation, George Mason University, Fairfax, VA.

Sutton, R.S. and A.G. Barto (1998) *Reinforcement Learning: An Introduction*, Cambridge: MIT Press.