

Lecture 7: Time Series Forecasting with Machine Learning Models

Isaiah Hull^{1,2}

¹BI Norwegian Business School

²CogniFrame

December 5, 2023



Lecture 7: Overview

1. Time Series Forecasting with Machine Learning.
2. Applications.

1. Time Series

Time Series

Time Series

- ▶ Empirical economics focuses on causal inference and hypothesis testing.
- ▶ Machine learning emphasizes prediction.
- ▶ Overlap in objectives when forecasting in economics and finance.

Time Series and Machine Learning

- ▶ Coulombe et al. (2019) discusses machine learning for time series econometrics.
- ▶ Potentially valuable tools: nonlinear models, regularization, cross validation, and alternative loss functions.
- ▶ Will focus on TensorFlow and deep learning models for time series.

Sequential Models in Machine Learning

- ▶ Specialized layers for neural networks to handle sequential data.
- ▶ Originally developed for natural language processing (NLP).
- ▶ Also applicable to time series contexts.
- ▶ Early example in Nakamura (2005) used neural networks for forecasting inflation.

Dense Neural Networks


- ▶ Discussed earlier.
- ▶ Not previously adapted for sequential data.
- ▶ Forecasting exercise: Predicting quarterly inflation similar to Nakamura (2005).

Forecasting Inflation

- ▶ Data on U.S. quarterly inflation from 1947:Q2 to 2020:Q2.
- ▶ Following Nakamura (2005), consider univariate models with lags of inflation.

Time Series

CPI Inflation: 1947:Q2 - 2020_Q2 (U.S. BLS)



Source: Hull (2021).

Pre-Processing Sequential Data


- ▶ Similar to text and image data, need to pre-process sequential data.
- ▶ Transform time series into fixed-length sequences.
- ▶ Decide on sequence length, e.g., number of lags for inputs.
- ▶ For a sequence length of three, use realizations in periods t , $t-1$, and $t-3$.

Illustrating Pre-Processing

- ▶ Split single time series into overlapping sequences of consecutive observations.
- ▶ Use sequences to predict inflation for a single quarter ahead.

Time Series

Partition of Time Series



Source: Hull (2021).

Sequence Generator for Inflation

```
# Import packages.
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.sequence import TimeseriesGenerator

# Load CSV data as pandas dataframe.
inflation = pd.read_csv('inflation.csv')

# Convert inflation column into numpy array.
inflation = np.array(inflation['Inflation'])

# Instantiate sequence generator.
generator = TimeseriesGenerator(inflation, inflation, length=4, batch_size=12)
```

Model Training

- ▶ Using the generator object for batches of data.
- ▶ Use Keras Sequential() model.
- ▶ Sequential API for layer stacking.
- ▶ Define input, hidden, and output layers.
- ▶ Compile model with mean squared error and adam optimizer.
- ▶ Use fit_generator() for training.

Train Neural Network with Sequences

Define sequential model.

```
model = tf.keras.models.Sequential()
```

Add input layer.

```
model.add(tf.keras.Input(shape=(4,)))
```

Define dense layer.

```
model.add(tf.keras.layers.Dense(2, activation="relu"))
```

Train Neural Network with Sequences

Define output layer.

```
model.add(tf.keras.layers.Dense(1, activation="linear"))
```

Compile the model.

```
model.compile(loss="mse", optimizer="adam")
```

Train the model.

```
model.fit_generator(generator, epochs=100)
```


Time Series

Train Neural Network with Sequences

#Train for 25 steps

Epoch 1/100

25/25 [=====] - loss: 4.3247

Epoch 100/100

25/25 [=====] - loss: 0.3816

Model Performance

- ▶ Between epochs 1-100: Mean squared error drops from 4.32 to 0.38.
- ▶ No regularization or test sample split.
- ▶ Possible overfitting, but model has only 13 trainable parameters.

Model Architecture Summary

```
print(model.summary())
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

dense_1 (Dense)	(None, 2)	10
-----------------	-----------	----

dense_1 (Dense)	(None, 1)	3
-----------------	-----------	---


Total params: 13

Trainable params: 13

Non-trainable params: 0

Time Series

True Values of Inflation vs. Model's Prediction

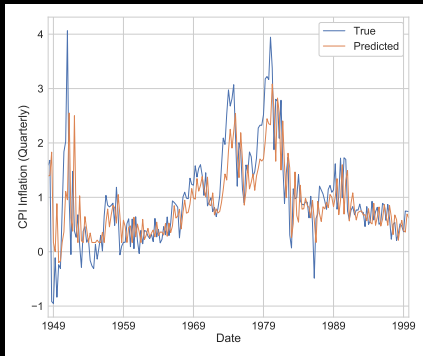


Source: Hull (2021).

Model Evaluation

- ▶ Use `model.predict_generator(generator)` for predictions.
- ▶ Examine overfitting.
- ▶ Use pre-2000 values for training and post-2000 for testing.
- ▶ Construct a separate generator for predictions.

Examining Potential Overfitting with Post-2000 Data



Source: Hull (2021).

Recurrent Neural Networks (RNN)

- ▶ RNN processes sequences using dense layers and recurrent layers.
- ▶ Inputs can be word vectors, musical notes, inflation measurements.
- ▶ Treatment follows Goodfellow et al. (2017).

Recurrent Layer Structure

- ▶ Recurrent layer consists of cells.
- ▶ Each cell takes input value $x(t)$ and state $h(t - 1)$.
- ▶ Produces an output value $o(t)$.

Multiplication in RNN Cell

$$a(t) = b + Wh(t - 1) + Ux(t)$$

- ▶ Take the state of the series, $h(t-1)$.
- ▶ Multiply by weights, W .
- ▶ Take the input value, $x(t)$, and multiply by a separate weights, U .
- ▶ Sum both terms together, along with a bias term, b .

Activation Function in RNN Cell

$$h(t) = \tanh(a(t))$$

- ▶ Take output of multiplication step.
- ▶ Pass output to hyperbolic tangent activation function.
- ▶ Output is the updated state of the system, $h(t)$.

Output from RNN Cell

$$o(t) = c + Vh(t)$$


- ▶ Multiply updated state by a separate set of weights, V .
- ▶ Add bias term, c .

Model Details

- ▶ In our example: inflation is the only feature.
- ▶ $x(t)$, W , U , and V are scalars.
- ▶ Weights are shared across time periods.
- ▶ Only five parameters needed for one RNN cell layer.

Time Series

RNN (top) and unrolled RNN cell (bottom)



Source: Hull (2021).

RNN Illustration

- ▶ Pink nodes: input values (lags of inflation).
- ▶ Orange node: inflation for the next quarter (target).
- ▶ Blue nodes: individual RNN cells in an RNN layer.

Unrolled RNN Cell

- ▶ State combined with input to get next state.
- ▶ Last step gives output $o(t)$.
- ▶ Output is input to a final dense layer for next quarter's inflation prediction.

RNN Characteristics

- ▶ Retains a state for sequential data.
- ▶ Updates state at each step.
- ▶ Reduces parameters through weight-sharing.
- ▶ No time-specific weights needed.
- ▶ Can handle sequences of varying length.

Model Definition

- ▶ Define RNN model.
- ▶ Model similar in complexity to previous dense network.
- ▶ Uses 'SimpleRNN' layer with 2 RNN cells.

Time Series

Sequence Generator for Inflation

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.sequence import TimeseriesGenerator

# Load data.
inflation = pd.read_csv(data_path+'inflation.csv')

# Convert to numpy array.
inflation = np.array(inflation['Inflation'])
```

Sequence Generator for Inflation

Add dimension.

```
inflation = np.expand_dims(inflation, 1)
```

Instantiate time series generator.

```
train_generator = TimeseriesGenerator(  
inflation[:211], inflation[:211],  
length = 4, batch_size = 12)
```

Define RNN Model in Keras

Define sequential model.

```
model = tf.keras.models.Sequential()
```

Define recurrent layer.

```
model.add(tf.keras.layers.SimpleRNN(2, input_shape=(4, 1)))
```

Define output layer.

```
model.add(tf.keras.layers.Dense(1, activation="linear"))
```

Define RNN Model in Keras

Compile the model.

```
model.compile(loss="mse", optimizer="adam")
```

Fit model to data using generator.


```
model.fit_generator(train_generator, epochs=100)
```

Model Compilation and Training

- ▶ Mean squared error (MSE) is lower compared to dense network.
- ▶ Test sample performance post-2000 remains stable.

Time Series

Test Sample Performance Post-2000



Source: Hull (2021).

Evaluate RNN Model in Keras

Epoch 1/100

18/18 [=====] - 1s 31ms/step - loss: 0.9206

Epoch 100/100

18/18 [=====] - 0s 2ms/step - loss: 0.2594

RNN vs. Dense Networks

- ▶ RNN has fewer parameters than dense networks.
- ▶ Single RNN cell layer requires 5 parameters.
- ▶ Examine architecture using 'model.summary()'.
- ▶ RNN model in this example: 11 parameters.
- ▶ Dense network used previously had more parameters.

RNN Architecture in Keras

Print model summary.

```
print(model.summary())
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

simple_rnn_1 (SimpleRNN)	(None, 2)	8
--------------------------	-----------	---

dense_1 (Dense)	(None, 1)	3
-----------------	-----------	---

Total params: 11

Trainable params: 11

Non-trainable params: 0

RNN in Practice

- ▶ Challenges with basic RNN:
 - ▶ Vanishing gradient problem.
 - ▶ Original RNN struggles with long data sequences.
 - ▶ Doesn't account for distant temporal relationships.

Long short-term memory (LSTM)

- ▶ RNNs suffer from vanishing gradient with long sequences.
- ▶ Solution: Use gated RNN cells.
- ▶ Two common gated RNN cells:
 - ▶ Long short-term memory (LSTM)
 - ▶ Gated recurrent units (GRUs)

Long short-term memory (LSTM)

- ▶ Introduced by Hochreiter and Schmidhuber (1997).
- ▶ Uses operators to limit flow of information in long sequences.

Operations in LSTM

- ▶ Following Goodfellow et al. (2017) for LSTM operations.
- ▶ Three gates in LSTM:
 - ▶ Forget gate.
 - ▶ External input gate.
 - ▶ Output gate.
- ▶ Gates control flow of information in LSTM cell.

Equations for LSTM Gates

- Forget gates:

$$f(t) = \sigma(b^f + W^f h(t-1) + U^f x(t))$$

- External input gates:

$$g(t) = \sigma(b^g + W^g h(t-1) + U^g x(t))$$

- Output gates:

$$q(t) = \sigma(b^q + W^q h(t-1) + U^q x(t))$$

Internal States of LSTM

- ▶ Each gate has unique weights and biases.
- ▶ Gating procedure can be learned.
- ▶ Internal states are updated using definitions provided.
- ▶ Incorporates forget gate, external input gate, input sequence, and state.

Equations for LSTM Gates

- ▶ Internal state:

$$s(t) = f^t s(t-1) + g(t) \sigma(b + Wh(t-1) + Ux(t))$$

- ▶ Hidden state:

$$h(t) = \tanh(s(t))q(t)$$

Use of Gates in LSTM

- ▶ Gates increase the number of model parameters.
- ▶ Significant improvement in handling long sequences.
- ▶ LSTM is typically the baseline ML model in time series analysis.

LSTM Model Definition and Training

- ▶ Use `tf.keras.layers.LSTM()` instead of `tf.keras.layers.SimpleRNN()`.
- ▶ Train for 100 epochs.
- ▶ Mean squared error is higher for LSTM than RNN after 100 epochs.
 - ▶ More weights in LSTM require more training epochs.
- ▶ Benefit from LSTM grows in sequence length.

Train an LSTM Model in Keras

Define sequential model.

```
model = tf.keras.models.Sequential()
```

Define recurrent layer.

```
model.add(tf.keras.layers.LSTM(2, input_shape=(4, 1)))
```

Define output layer.

```
model.add(tf.keras.layers.Dense(1, activation="linear"))
```

Train an LSTM Model in Keras

Compile the model.

```
model.compile(loss="mse", optimizer="adam")
```

Train the model.

```
model.fit_generator(train_generator, epochs=100)
```

Time Series

Train an LSTM Model in Keras

Epoch 1/100

18/18 [=====] - 1s 62ms/step - loss: 3.1697

Epoch 100/100

18/18 [=====] - 0s 3ms/step - loss: 0.5873

LSTM Model's Architecture

- ▶ LSTM cells introduce additional operations:
 - ▶ Forget gate
 - ▶ External input gate
 - ▶ Output gate
- ▶ Each gate requires its own set of parameters.
- ▶ LSTM layer uses 32 parameters, which is 4 times as many as the RNN.

Summarize LSTM Architecture in a Keras Model

Print model architecture.

```
print(model.summary())
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

lstm_1 (LSTM)	(None, 2)	32
---------------	-----------	----

dense_1 (Dense)	(None, 1)	3
-----------------	-----------	---

Total params: 35

Trainable params: 35

Non-trainable params: 0

Intermediate Hidden States

- ▶ By convention, LSTM only uses the final value of the hidden state.
- ▶ For instance:
 - ▶ Model uses $h(t)$.
 - ▶ Does not use $h(t-1)$, $h(t-2)$, $h(t-3)$.

Use of Intermediate Hidden States

- ▶ Recent research shows benefits from using intermediate hidden states.
- ▶ Especially effective for modeling long-term dependencies in NLP.
- ▶ Typically done within attention models.

Incorporating Hidden States in LSTM

- ▶ Modify LSTM cells to return hidden states.
- ▶ Set `return_sequences` to `True`.

Code Implementation

Modify LSTM to return hidden states

```
LSTM(return_sequences=True)
```

- ▶ After modification, we can review model's architecture using the summary() method.

Time Series

Code Implementation

Define sequential model.

```
model = tf.keras.models.Sequential()
```

Define recurrent layer to return hidden states.

```
model.add(tf.keras.layers.LSTM(2, return_sequences=True, input_shape=(4, 1)))
```

Define output layer.

```
model.add(tf.keras.layers.Dense(1, activation="linear"))
```

Summarize model architecture.

```
model.summary()
```

Code Implementation

Layer (type)	Output Shape	Param #
--------------	--------------	---------

lstm_1 (LSTM)	(None, 4, 2)	32
---------------	--------------	----

dense_1 (Dense)	(None, 4, 1)	3
-----------------	--------------	---

Total params: 35

Trainable params: 35

Non-trainable params: 0

Unusual Model Architecture

- ▶ Model outputs a 4x1 vector, not a scalar prediction for each observation.
- ▶ This is due to the LSTM layer:
 - ▶ Outputs 4x1 vectors from its two LSTM cells.

Making Use of LSTM Output

- ▶ Several methods to incorporate the LSTM output.
- ▶ Example: **Stacked LSTM** (Graves et al., 2013).
 - ▶ Pass full sequence hidden states to a second LSTM layer.
 - ▶ Adds depth to network for multiple levels of representation.

Stacked LSTM Implementation

- ▶ Define a model with two LSTM layers.
- ▶ First LSTM layer:
 - ▶ Three LSTM cells.
 - ▶ Input shape of (4,1).
 - ▶ Set return_sequences to True.
 - ▶ Outputs 4x1 sequence of hidden states.

Second LSTM Layer

- ▶ Accepts 3-tensor (4x1x3) from the first LSTM layer.
- ▶ Contains two cells.
- ▶ Only returns the final hidden states.
- ▶ Does not return intermediate state values.

Time Series

Define a Stacked LSTM Model

Define sequential model.

```
model = tf.keras.models.Sequential()
```

Define recurrent layer to return hidden states.

```
model.add(tf.keras.layers.LSTM(3, return_sequences=True, input_shape=(4, 1)))
```

Define second recurrent layer.

```
model.add(tf.keras.layers.LSTM(2))
```

Define output layer.

```
model.add(tf.keras.layers.Dense(1, activation="linear"))
```

Summarize Stacked LSTM Architecture

Summarize model architecture.

model.summary()

Layer (type)	Output Shape	Param #
=====		
lstm_1 (LSTM)	(None, 4, 3)	60
=====		
lstm_2 (LSTM)	(None, 2)	48
=====		
dense_1 (Dense)	(None, 1)	3
=====		
Total params: 111		
Trainable params: 111		
Non-trainable params: 0		

Time Series

Focus and Context

- ▶ Apply forecasting methods.
- ▶ Based examples on univariate inflation forecasting (Nakamura, 2005).
- ▶ All methods applicable to multivariate settings.

Multivariate Forecasting

- ▶ Use both LSTM model and Gradient Boosted Trees.
- ▶ Forecast inflation monthly.
- ▶ Use five features instead of one.

Data Preview and Model Implementation

- ▶ Prepare data and define model.
- ▶ Perform multivariate forecast using:
 - ▶ LSTM
 - ▶ Gradient boosted trees

Time Series

Features Details

- ▶ New features added:
 - ▶ Unemployment (measured in first differences).
 - ▶ Hours worked in manufacturing.
 - ▶ Hourly earnings in manufacturing.
 - ▶ Money supply measure (M1).
- ▶ Level variables transformed using percentage changes from previous period.

Time Series

Load and Preview Inflation Forecast Data

```
import pandas as pd
```

```
# Load data.
```

```
macroData = pd.read_csv(data_path+'macrodata.csv', index_col = 'Date')
```

```
# Preview data.
```

```
print(macroData.round(1).tail())
```

Date	Inflation	Unemp	Hours	Earnings	M1
12/1/19	-0.1	0.1	0.5	0.2	0.7
01/1/20	0.4	0.6	-1.7	-0.1	0.0
02/1/20	0.3	-0.2	0.0	0.4	0.8
03/1/20	-0.2	0.8	-0.2	0.4	6.4
04/1/20	-0.7	9.8	-6.8	0.5	12.9

LSTM Overview

- ▶ Recall data preparation for LSTM:
 - ▶ Instantiate a generator.
 - ▶ Convert target and features to `np.array()` objects.

Data Generators

- ▶ Two data generators:
 - ▶ Training data
 - ▶ Test data

Time Series

Sequence Lengths

- ▶ Previous setup:
 - ▶ Quarterly data
 - ▶ 4-quarter sequence lengths
- ▶ Current setup:
 - ▶ Monthly data
 - ▶ 12-month sequence lengths

Time Series

Prepare Data for Use in LSTM Model

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing.sequence import TimeseriesGenerator

# Define target and features.
target = np.array(macroData['Inflation'])
features = np.array(macroData)

# Define train generator.
train_generator = TimeseriesGenerator(features[:393],
                                       target[:393], length = 12, batch_size = 6)
```

Prepare Data for Use in LSTM Model

Define test generator.

```
test_generator = TimeseriesGenerator(features[393:],  
                                     target[393:], length = 12, batch_size = 6)
```

Generators Defined

- ▶ Train model.
- ▶ Model Specifications:
 - ▶ Two LSTM cells.
 - ▶ Input shape adjusted:
 - ▶ 12 elements in each sequence.
 - ▶ Five features.

Training Outcome

- ▶ Concerns in macroeconomic forecasting:
 - ▶ Longer sequences and more variables.
 - ▶ Typically, concerns about number of parameters.

Model Parameters

- ▶ Despite longer sequence length:
 - ▶ Doesn't increase the number of parameters.
- ▶ Model's total parameters:
 - ▶ Only 67 parameters.

Time Series

Load and Preview Inflation Forecast Data

Define sequential model.

```
model = tf.keras.models.Sequential()
```

Define LSTM model with two cells.

```
model.add(tf.keras.layers.LSTM(2, input_shape=(12, 5)))
```

Define output layer.

```
model.add(tf.keras.layers.Dense(1, activation="linear"))
```

Compile the model.

```
model.compile(loss="mse", optimizer="adam")
```

Time Series

Load and Preview Inflation Forecast Data

Train the model.

```
model.fit_generator(train_generator, epochs=100)
```

Epoch 1/20

64/64 [=====] - 2s 26ms/step - loss: 0.3065

Epoch 20/20

64/64 [=====] - 0s 6ms/step - loss: 0.0663

Use MSE to Evaluate Train and Test Sets

Evaluate training set using MSE.

```
model.evaluate_generator(train_generator)
```

0.06527029448989197

Evaluate test set using MSE.

```
model.evaluate_generator(test_generator)
```

0.15478561431742632

Model Evaluation

- ▶ Compare training sample results to test sample results.
- ▶ Training set performance generally better than test set.
- ▶ Not uncommon, but indicative of overfitting.

Addressing Performance Gap

- ▶ If gap between training and test set performance is large:
 - ▶ Consider regularization.
 - ▶ Consider terminating training earlier.
 - ▶ Reduce number of epochs.

Gradient Boosted Trees

- ▶ Compare GBT to deep learning.

Data Preparation for Gradient Boosting

- ▶ Similar to LSTMs, data splitting into sequences is required.
- ▶ In TensorFlow, GBT trained using Estimator API.
- ▶ Must define feature columns for each of the five features.

Data Generation Functions

- ▶ Define functions to generate data for both training and testing.
- ▶ Evaluate overfitting, analogous to the LSTM example.
- ▶ Sample split:
 - ▶ Train set: years before 2000.
 - ▶ Test set: years post-2000.

Use MSE to Evaluate Train and Test Sets

Define lagged inflation feature column.

```
inflation = tf.feature_column.numeric_column("inflation")
```

Define unemployment feature column.

```
unemployment = tf.feature_column.numeric_column("unemployment")
```

Define hours feature column.

```
hours = tf.feature_column.numeric_column("hours")
```

Time Series

Use MSE to Evaluate Train and Test Sets

Define earnings feature column.

```
earnings = tf.feature_column.numeric_column("earnings")
```

Define M1 feature column.

```
m1 = tf.feature_column.numeric_column("m1")
```

Define feature list.

```
feature_list = [inflation, unemployment, hours, earnings, m1]
```

Time Series

Define the Data Generation Functions

Define input function for training data.

```
def train_data():  
    train = macroData.iloc[:392]  
    features = {"inflation": train["Inflation"],  
               "unemployment": train["Unemployment"],  
               "hours": train["Hours"],  
               "earnings": train["Earnings"],  
               "m1": train["M1"]}  
    labels = macroData["Inflation"].iloc[1:393]  
    return features, labels
```

Time Series

Define the Data Generation Functions

Define input function for test data.

```
def test_data():  
    test = macroData.iloc[393:-1]  
    features = {"inflation": test["Inflation"],  
               "unemployment": test["Unemployment"],  
               "hours": test["Hours"],  
               "earnings": test["Earnings"],  
               "m1": test["M1"]}  
    labels = macroData["Inflation"].iloc[394:]  
    return features, labels
```

Time Series

Train and Evaluate Model

Instantiate boosted trees regressor.

```
model = tf.estimator.BoostedTreesRegressor(feature_columns =  
    feature_list, n_batches_per_layer = 1)
```

Train model.

```
model.train(train_data, steps=100)
```

Evaluate train and test set.

```
train_eval = model.evaluate(train_data, steps = 1)
```

```
test_eval = model.evaluate(test_data, steps = 1)
```

Print results.

```
print(pd.Series(train_eval))
```

```
print(pd.Series(test_eval))
```

Time Series

Train and Evaluate Model

average_loss	0.010534
label/mean	0.416240
loss	0.010534
prediction/mean	0.416263
global_step	100.000000
dtype: float64	

average_loss	0.145123
label/mean	0.172864
loss	0.145123
prediction/mean	0.286285
global_step	100.000000
dtype: float64	

ML in Economics and Finance

- ▶ ML focuses on prediction.
- ▶ Economics and finance often aim for causal inference and hypothesis testing.
- ▶ Overlap in fields: forecasting.

Time Series

Time Series Forecasting with ML

- ▶ Emphasis on deep learning models.
- ▶ Gradient boosted trees in TensorFlow.
- ▶ Early use of neural network in economics for forecasting: Nakamura (2005).
- ▶ Modern models: RNNs, LSTMs, stacked LSTMs.
- ▶ Applications extend to areas like NLP.

Further Reading

- ▶ Macroeconomic time series forecasting: Cook and Hall (2017).
- ▶ Stock return & bond premium forecasting: Heaton et al. (2016), Messmer (2017), Rossi (2018), and Chen et al. (2019).
- ▶ High-dimensional time series regression and nowcasting: Babii et al. (2019, 2020).

2. Applications

Applications

Colab Tutorial

- ▶ Time Series Prediction

References I

- Babii, A., E. Ghysels, and J. Striaukas (2019) "Inference for High-Dimensional Regressions with Heteroskedasticity and Auto-correlation," *arXiv preprint*.
- (2020) "Machine Learning Time Series Regressions with an Application to Nowcasting," *arXiv preprint*.
- Chen, L., M. Pelger, and J. Zhu (2019) "Deep Learning in Asset Pricing," *arXiv preprint*.
- Cook, T.R. and A.S. Hall (2017) "Macroeconomic Indicator Forecasting with Deep Neural Networks," Research Working Paper 17-11, Federal Reserve Bank of Kansas City.

References II

- Coulombe, P.G., M. Leroux, D. Stevanovic, and S. Surprenant (2019) "How is Machine Learning Useful for Macroeconomic Forecasting?" working papers, CIRANO.
- Goodfellow, I., Y. Bengio, and A. Courville (2017) *Deep Learning*, Cambridge, MA: MIT Press.
- Graves, A., A.-r. Mohamed, and G. Hinton (2013) "Speech Recognition with Deep Recurrent Neural Networks," *arXiv preprint*.
- Heaton, J.B., N.G. Polson, and J.H. Witte (2016) "Deep Learning for Finance: Deep Portfolios," *Applied Stochastic Models in Business and Industry*, 33 (1), 3–12.
- Hochreiter, S. and J. Schmidhuber (1997) "Long Short-term Memory," *Neural Computation*, 9 (8), 1735–1780.

References III

Hull, Isaiah (2021) *Machine Learning for Economics and Finance in TensorFlow 2*: Apress, 10.1007/978-1-4842-6373-0.

Messmer, M. (2017) “Deep Learning and the Cross-Section of Expected Returns,” ssrn working paper.

Nakamura, Emi (2005) “Inflation Forecasting Using a Neural Network,” *Economics Letters*, 85, 373–378.

Rossi, A.G. (2018) “Predicting Stock Market Returns with Machine Learning,” working paper.